


Dropping SFT was the move.

Four supervised runs, four regressions. The recipe we kept uses pure prompt engineering — 0 training records.

autoresearch · progress-update · DeepMath 500

159 experiments · 26 final-recipe samples



Prompt-only beats every fine-tuned variant we tried. Even 20 records was enough to collapse cadence. Over 200 A3B · 0 records · 95% CI [0.871, 0.880]